



even if you might not need it for your documents. We'll discuss this more in class.

5. Parsing exceptions should be handled with a framework-specific exception class (inherited from Exception).
6. The library will support three visualizations, one of which is very specific, the other two are left at your discretion.
  - a. (Specific) Text-to-Word Sankey diagram. Given the loaded texts and either a set of user-defined words OR the set of words drawn from the  $k$  most common words of each text file, generate a Sankey diagram from text name to word, where the thickness of the connection represents the wordcount of that word in the specified text.
  - b. (Flexible) Any type of visualization containing sub-plots, one *sub-plot* for each text file. For example, an array of word clouds, one for each text file would satisfy this requirement, but be sure to write your code so that the subplot array dimensions can handle from 2-10 submitted documents. If you want the dimensions of the subplot array could be overridden by the user, but you should define reasonable defaults.
  - c. (Flexible) Any type of comparative visualization that overlays information from each text file onto a *single* visualization.

## Library Architecture

I recommend that your library be given a useful name. (Imagine trying to share your library on GitHub but you called it "hw2" – who's going to want to download that?)

At minimum, you should implement the following methods:

- **load\_text(self, filename, label="", parser=None)**  
# Register a text file with the library. The label is an optional label you'll use in your  
# visualizations to identify the text
- **load\_stop\_words(stopfile)**  
# A list of common or stop words. These get filtered from each file automatically
- **wordcount\_sankey(self, word\_list=None, k=5)**  
# Map each text to words using a Sankey diagram, where the thickness of the line  
# is the number of times that word occurs in the text. Users can specify a particular  
# set of words, or the words can be the union of the  $k$  most common words across

# each text file (excluding stop words).

- ***your\_second\_visualization(self, misc\_parameters)***  
# What you do here is up to you, but it should produce a single visualization with one  
# subplot for each text file. Rendering subplots is a good, advanced skill to know!
- ***your\_third\_visualization(self, misc\_parameters)***  
# This is also left up to you, but it should be a single visualization that overlays data from  
# each of the text files. Make sure your visualization distinguishes the data from each  
text file using labels or a legend.

As you pre-process each text file, you'll want to store intermediate results: word counts, statistics, maybe the clean version of text, etc. A standard way to do this might be to have a state variable: **self.data** that is a dictionary. Presented below is a generic solution using dictionaries within dictionaries. The layout and content are ultimately up to you but try to come up with an approach that makes it easy to add custom parsers that store other kinds of information and custom visualizations that can access this information in standard ways.

```
{
    word_count: {text1: wordcount1, text2: wordcount2, .... }
    word_length: {text1, wordlength1, text2:wordlength2, ... }
    sentiment: {text1, sentiment1, text2:sentiment2, ... }
    .
    .
    .
}
```

## Data Sources

Identify 5-10 related documents. The type of file you choose is entirely up to you, but I recommend something different than presidential inaugural speeches. (Been there, done that.) Some possible ideas or data sources:

- Project Gutenberg ([www.gutenberg.org](http://www.gutenberg.org))
- Political speeches
- National constitutions
- Tweet compilations
- Corporate filings
- Letters/Journals/Diaries, for example:  
<https://content.lib.washington.edu/civilwarweb/index.html>

(It would be interesting to compare letters from soldiers fighting for the North vs South!)

- News articles or blog posts
- Religious texts
- Philosophical tracts
- Song lyrics

## Assignment Guidelines

Your submission should include a 2 to 3-page report in **PDF** format summarizing your data sources, insights, and conclusions.

You may work in groups of 1-4. For group submissions, one submission per group is sufficient. Please document the contributions of each team member to the code base in your final report. The submitter should identify each member of the group to ensure each team member receives credit.

Code for this *mini project* must be managed on the NEU Enterprise Git Server. Grant the instructors access or make the project public. Even though the code will live in a repo, please also submit your code files and project report to Gradescope so that the TAs can comment on specific lines of code.

We'll have in-class mini-presentations to share our work. Each group should use powerpoint to prepare a ONE SLIDE POSTER (PDF version to be submitted to GradeScope). Multi-slide posters will not be accepted.

## What to Submit

- Code (.py)
- 2 to 3-page report, including author contributions, visualizations, commentary and a link to your code repository (.pdf)
- 1 slide poster for the in-class mini-presentation (.pdf)